

# QBS181\_ProblemSet3

Matthew Chan

10/28/2021

## Part One

### Question One

Recall the HigherMe dataset. On the canvas homepage, I've linked a csv file which contains the data cleaning up to the step where we need to convert those invoices which are quarterly to monthly invoices. Using the same logic you applied in the excel project, convert quarterly invoices to monthly invoices in R. Display the first 10 rows of your updated dataset.

```
#your code here
```

```
library(dplyr)
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
## filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
## intersect, setdiff, setequal, union
```

```
library(lubridate)
```

```
##
```

```
## Attaching package: 'lubridate'
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
## date, intersect, setdiff, union
```

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.1 --
```

```
## v ggplot2 3.3.5      v purrr 0.3.4
```

```
## v tibble 3.1.4       v stringr 1.4.0
```

```
## v tidyr 1.1.3        v forcats 0.5.1
```

```
## v readr 2.0.1
```

```
## -- Conflicts ----- tidyverse_conflicts() --
## x lubridate::as.difftime() masks base::as.difftime()
## x lubridate::date() masks base::date()
## x dplyr::filter() masks stats::filter()
## x lubridate::intersect() masks base::intersect()
## x dplyr::lag() masks stats::lag()
## x lubridate::setdiff() masks base::setdiff()
## x lubridate::union() masks base::union()
```

```
library(data.table)
```

```
##
```

```
## Attaching package: 'data.table'
```

```
## The following object is masked from 'package:purrr':
```

```
##
```

```
## transpose
```

```
## The following objects are masked from 'package:lubridate':
```

```
##
```

```
## hour, isoweek, mday, minute, month, quarter, second, wday, week,
```

```
## yday, year
```

```
## The following objects are masked from 'package:dplyr':
```

```
##
```

```
## between, first, last
```

```
Higherme <- read.csv("HigherMeDataForRData.csv")
```

```
names(Higherme)[names(Higherme) == "i..Invoice.Number"] <- "Invoice_Number"
Higherme$Invoice.Date <- as.Date(Higherme$Invoice.Date, format = "%m/%d/%y")
head(Higherme)
```

```
## Invoice_Number Invoice.Date Amount Status Paid.On Refunded.Amount
## 1 49236 2021-09-14 $115.00 Paid 9/14/21 22:33 0
## 2 49235 2021-09-14 $115.00 Paid 9/14/21 22:33 0
## 3 49213 2021-09-14 $360.00 Paid 9/14/21 21:42 0
## 4 49212 2021-09-14 $4,455.00 Paid 9/14/21 21:28 0
## 5 49183 2021-09-14 $57.50 Paid 9/14/21 5:05 0
## 6 49149 2021-09-14 $2,640.00 Paid 9/14/21 0:02 0
## Recurring Quarterly.Monthly First.Invoice Company Tax.Total Amount.Due
## 1 FALSE MONTHLY FALSE Company 1 $- $-
## 2 FALSE MONTHLY FALSE Company 1 $- $-
## 3 TRUE QUARTERLY FALSE Company 2 $- $-
## 4 TRUE QUARTERLY FALSE Company 3 $- $-
## 5 FALSE MONTHLY FALSE Company 3 $- $-
## 6 TRUE QUARTERLY FALSE Company 4 $- $-
## Adjustments Credits.Applied Payments Write.Off.Amount Currency
## 1 $- $- $115.00 $- USD
## 2 $- $- $115.00 $- USD
## 3 $- $- $360.00 $- USD
## 4 $- $- $4,455.00 $- USD
```

```
## 5      $-      $-      $57.50      $-      USD
## 6      $-      $-      $2,640.00    $-      USD
##      Due.Date Customer.Billing.Country X
## 1 9/14/21 22:33      NA
## 2 9/14/21 22:33      NA
## 3 9/14/21 21:42      US NA
## 4 9/14/21 21:28      NA
## 5 9/14/21 5:05      NA
## 6 9/14/21 0:02      NA
```

*# Remove dollar signs function as provided in class*

```
rmCurrency<-function(x){

  x<-trimws(x) #trim whitespace

  if(grepl("\\$",x[1])){ #if '$' found in x
    x<-sub("\\$", "",x) #remove $
    x<-sub("\\,", "",x) #remove ,
    x[x=="-"]<-0 #recode zeros
  }
  return(x)
}

Higherme$Amount <- rmCurrency(Higherme$Amount)
Higherme$Amount <- as.integer(Higherme$Amount)
Higherme$Amount <- Higherme$Amount/ 3
head(Higherme)
```

```
## Invoice_Number Invoice.Date Amount Status Paid.On Refunded.Amount
## 1 49236 2021-09-14 38.33333 Paid 9/14/21 22:33 0
## 2 49235 2021-09-14 38.33333 Paid 9/14/21 22:33 0
## 3 49213 2021-09-14 120.00000 Paid 9/14/21 21:42 0
## 4 49212 2021-09-14 1485.00000 Paid 9/14/21 21:28 0
## 5 49183 2021-09-14 19.00000 Paid 9/14/21 5:05 0
## 6 49149 2021-09-14 880.00000 Paid 9/14/21 0:02 0
## Recurring Quarterly.Monthly First.Invoice Company Tax.Total Amount.Due
## 1 FALSE MONTHLY FALSE Company 1 $- $-
## 2 FALSE MONTHLY FALSE Company 1 $- $-
## 3 TRUE QUARTERLY FALSE Company 2 $- $-
## 4 TRUE QUARTERLY FALSE Company 3 $- $-
## 5 FALSE MONTHLY FALSE Company 3 $- $-
## 6 TRUE QUARTERLY FALSE Company 4 $- $-
## Adjustments Credits.Applied Payments Write.Off.Amount Currency
## 1 $- $- $115.00 $- USD
## 2 $- $- $115.00 $- USD
## 3 $- $- $360.00 $- USD
## 4 $- $- $4,455.00 $- USD
## 5 $- $- $57.50 $- USD
## 6 $- $- $2,640.00 $- USD
## Due.Date Customer.Billing.Country X
## 1 9/14/21 22:33 NA
## 2 9/14/21 22:33 NA
## 3 9/14/21 21:42 US NA
## 4 9/14/21 21:28 NA
## 5 9/14/21 5:05 NA
```

## 6 9/14/21 0:02

NA

```
quarterly2 <- Higherme %>%  
  filter(Quarterly.Monthly == "QUARTERLY") %>%  
  arrange(Invoice_Number)  
quarterly3 <- Higherme %>%  
  filter(Quarterly.Monthly == "QUARTERLY") %>%  
  arrange(Invoice_Number)
```

```
quarterly2$Invoice.Date <- ymd(quarterly2$Invoice.Date) - 30  
head(quarterly2)
```

##	Invoice_Number	Invoice.Date	Amount	Status	Paid.On	Refunded.Amount
## 1	5937	2018-02-12	70.0000	Paid	3/14/18 21:40	0
## 2	6406	2018-03-15	70.0000	Paid	4/14/18 21:40	0
## 3	6701	2018-04-03	25.0000	Paid	5/3/18 15:53	0
## 4	6714	2018-04-04	883.3333	Paid	5/4/18 13:01	0
## 5	6766	2018-04-08	166.6667	Paid	5/8/18 17:06	300
## 6	6788	2018-04-09	704.3333	Paid	5/9/18 20:26	0
##	Recurring	Quarterly.Monthly	First.Invoice	Company	Tax.Total	Amount.Due
## 1	TRUE	QUARTERLY	TRUE	Company 2	\$-	\$-
## 2	TRUE	QUARTERLY	FALSE	Company 2	\$-	\$-
## 3	TRUE	QUARTERLY	TRUE	Company 28	\$-	\$-
## 4	TRUE	QUARTERLY	TRUE	Company 27	\$-	\$-
## 5	TRUE	QUARTERLY	TRUE	Company 19	\$-	\$-
## 6	TRUE	QUARTERLY	FALSE	Company 27	\$-	\$-
##	Adjustments	Credits.Applied	Payments	Write.Off.Amount	Currency	
## 1	\$-	\$-	\$210.00	\$-	USD	
## 2	\$-	\$-	\$210.00	\$-	USD	
## 3	\$-	\$-	\$75.00	\$-	USD	
## 4	\$-	\$-	\$2,650.00	\$-	USD	
## 5	\$-	\$-	\$500.00	\$-	CAD	
## 6	\$-	\$-	\$2,113.44	\$-	USD	
##	Due.Date	Customer.Billing.Country	X			
## 1	3/14/18 21:40	US	NA			
## 2	4/14/18 21:40	US	NA			
## 3	5/3/18 15:53	NA				
## 4	5/4/18 13:01	NA				
## 5	5/8/18 17:06	CA	NA			
## 6	5/9/18 20:26	NA				

```
quarterly3$Invoice.Date <- ymd(quarterly3$Invoice.Date) - 60  
head(quarterly3)
```

##	Invoice_Number	Invoice.Date	Amount	Status	Paid.On	Refunded.Amount
## 1	5937	2018-01-13	70.0000	Paid	3/14/18 21:40	0
## 2	6406	2018-02-13	70.0000	Paid	4/14/18 21:40	0
## 3	6701	2018-03-04	25.0000	Paid	5/3/18 15:53	0
## 4	6714	2018-03-05	883.3333	Paid	5/4/18 13:01	0
## 5	6766	2018-03-09	166.6667	Paid	5/8/18 17:06	300
## 6	6788	2018-03-10	704.3333	Paid	5/9/18 20:26	0
##	Recurring	Quarterly.Monthly	First.Invoice	Company	Tax.Total	Amount.Due

```
## 1      TRUE      QUARTERLY      TRUE Company 2      $-      $-
## 2      TRUE      QUARTERLY      FALSE Company 2      $-      $-
## 3      TRUE      QUARTERLY      TRUE Company 28      $-      $-
## 4      TRUE      QUARTERLY      TRUE Company 27      $-      $-
## 5      TRUE      QUARTERLY      TRUE Company 19      $-      $-
## 6      TRUE      QUARTERLY      FALSE Company 27      $-      $-
##      Adjustments Credits.Applied      Payments Write.Off.Amount      Currency
## 1      $-      $-      $210.00      $-      USD
## 2      $-      $-      $210.00      $-      USD
## 3      $-      $-      $75.00      $-      USD
## 4      $-      $-      $2,650.00      $-      USD
## 5      $-      $-      $500.00      $-      CAD
## 6      $-      $-      $2,113.44      $-      USD
##      Due.Date Customer.Billing.Country      X
## 1 3/14/18 21:40      US NA
## 2 4/14/18 21:40      US NA
## 3 5/3/18 15:53      NA
## 4 5/4/18 13:01      NA
## 5 5/8/18 17:06      CA NA
## 6 5/9/18 20:26      NA
```

```
# Binding my quarterly3 and quarterly2 dataframe together
merged <- rbind(quarterly3, quarterly2)
```

```
# Putting together the final dataframe
merged_final <- rbind(merged, Higherme)
```

```
# Changing all the values to Monthly from Quarterly
merged_final$Quarterly.Monthly[merged_final$Quarterly.Monthly=="QUARTERLY"]<- "MONTHLY"
merged_final <- merged_final %>%
  arrange(merged_final$Invoice_Number)
head(merged_final)
```

```
##      Invoice_Number Invoice.Date Amount Status      Paid.On Refunded.Amount
## 1      5937      2018-01-13      70      Paid 3/14/18 21:40      0
## 2      5937      2018-02-12      70      Paid 3/14/18 21:40      0
## 3      5937      2018-03-14      70      Paid 3/14/18 21:40      0
## 4      6406      2018-02-13      70      Paid 4/14/18 21:40      0
## 5      6406      2018-03-15      70      Paid 4/14/18 21:40      0
## 6      6406      2018-04-14      70      Paid 4/14/18 21:40      0
##      Recurring Quarterly.Monthly First.Invoice      Company Tax.Total Amount.Due
## 1      TRUE      MONTHLY      TRUE Company 2      $-      $-
## 2      TRUE      MONTHLY      TRUE Company 2      $-      $-
## 3      TRUE      MONTHLY      TRUE Company 2      $-      $-
## 4      TRUE      MONTHLY      FALSE Company 2      $-      $-
## 5      TRUE      MONTHLY      FALSE Company 2      $-      $-
## 6      TRUE      MONTHLY      FALSE Company 2      $-      $-
##      Adjustments Credits.Applied      Payments Write.Off.Amount      Currency      Due.Date
## 1      $-      $-      $210.00      $-      USD 3/14/18 21:40
## 2      $-      $-      $210.00      $-      USD 3/14/18 21:40
## 3      $-      $-      $210.00      $-      USD 3/14/18 21:40
## 4      $-      $-      $210.00      $-      USD 4/14/18 21:40
## 5      $-      $-      $210.00      $-      USD 4/14/18 21:40
```

```
## 6      $-      $-      $210.00      $-      USD 4/14/18 21:40
## Customer.Billing.Country X
## 1      US NA
## 2      US NA
## 3      US NA
## 4      US NA
## 5      US NA
## 6      US NA
```

## Question Two

Recall the hospital database. Recreate the physician-referral table from Question 9 on your SQL homework using R. I've loaded a zip folder of all the necessary csv files to do this on canvas, as well as a snapshot of what the table should look like for your reference.

```
#your code here
physician <- fread("Physician.csv")
patient <- fread("Patient.csv")
appointment <- fread("Appointment.csv")
undergoes <- fread("Undergoes.csv")
med_procedure <- fread("Medical_Procedure.csv")
affiliated_with <- fread("Affiliated_With.csv")
department <- fread("Department.csv")
```

```
sub_patient <- patient %>%
  left_join(appointment, by=c("SSN" = "Patient")) %>%
  filter(PCP!= Physician) %>%
  mutate("Patient_SSN" = SSN, "Cost" = 0)
```

```
Table2_join <- patient %>% left_join(undergoes, by=c("SSN"="Patient")) %>% left_join(med_procedure, by=
```

```
# Merging both into the dataframe
rawdata <- union_all(sub_patient, Table2_join )
rawdata
```

```
##      SSN      Name      Address      Phone InsuranceID PCP
## 1: 100000001 John Smith 42 Foobar Lane 555-0256 68476213 1
## 2: 100000004 Dennis Doe 1100 Foobaz Avenue 555-2048 68421879 3
## 3: 100000004 Dennis Doe 1100 Foobaz Avenue 555-2048 68421879 3
## 4: 100000004 Dennis Doe 1100 Foobaz Avenue 555-2048 68421879 3
## 5: 100000006 Rebecca Carrannante 3 Orange Court 555-7601 46268734 1
## ---
## 181: NA <NA> <NA> <NA> NA 1
## 182: NA <NA> <NA> <NA> NA 2
## 183: NA <NA> <NA> <NA> NA 1
## 184: NA <NA> <NA> <NA> NA 1
## 185: NA <NA> <NA> <NA> NA 1
## AppointmentID PrepNurse Physician      Start      End
## 1: 76983231 NA 3 2008-04-26 12:00:00 2008-04-26 13:00:00
## 2: 46846589 103 4 2008-04-25 10:00:00 2008-04-25 11:00:00
## 3: 59871321 NA 4 2008-04-26 10:00:00 2008-04-26 11:00:00
## 4: 86213939 102 9 2008-04-27 10:00:00 2008-04-21 11:00:00
```

```
## 5:      33556494      101      9 2008-05-12 00:00:00 2008-05-12 00:32:00
## ---
## 181:      NA      NA      7      <NA>      <NA>
## 182:      NA      NA      6      <NA>      <NA>
## 183:      NA      NA      7      <NA>      <NA>
## 184:      NA      NA      7      <NA>      <NA>
## 185:      NA      NA      3      <NA>      <NA>
##      ExaminationRoom Patient_SSN Cost
## 1:      C      100000001      0
## 2:      B      100000004      0
## 3:      C      100000004      0
## 4:      A      100000004      0
## 5:      B      100000006      0
## ---
## 181:      <NA>      100000100 3750
## 182:      <NA>      100000101 3750
## 183:      <NA>      100000102 1500
## 184:      <NA>      100000103 1500
## 185:      <NA>      100000103 1500
```

```
# Making the department summary table by merging the affiliated_with table
```

```
department_summary <- merge(affiliated_with, department, by.x = 'Department', by.y='DepartmentID')
department_summary <- department_summary %>%
  select(Department, Name) %>%
  distinct(Department, Name) %>%
  rename(Department_Name = Name)
```

```
# Making EmployeeID-Physician name as referential keys
```

```
physician_summary <- merge(affiliated_with, physician, by.x = 'Physician', by.y='EmployeeID')
physician_summary <- physician_summary %>% select(Physician, Department, Name) %>% rename(Physician_Name = Physician)
```

```
# Make a big reference table with the Physician and Department summaries combined
```

```
physician_department_summary <- merge(physician_summary, department_summary, by='Department')
physician_department_summary
```

```
##      Department Physician      Physician_Name      Department_Name
## 1:      1      1      John Dorian General Medicine
## 2:      1      2      Elliot Reid General Medicine
## 3:      1      3      Christopher Turk General Medicine
## 4:      1      4      Percival Cox General Medicine
## 5:      1      5      Bob Kelso General Medicine
## 6:      1      7      John Wen General Medicine
## 7:      1      8      Keith Dudemeister General Medicine
## 8:      2      3      Christopher Turk      Surgery
## 9:      2      6      Todd Quinlan      Surgery
## 10:      2      7      John Wen      Surgery
## 11:      3      9      Molly Clock      Psychiatry
```

```
# I replace the column name Referral and Referring Physicians in this table
```

```
finaltable <- rawdata %>%
  left_join(physician_department_summary, by=c("PCP" = "Physician")) %>%
  rename('Referring_Physician' = Physician_Name) %>%
```

```

left_join(physician_department_summary, by=c("Physician")) %>%
mutate(Referral = Physician_Name) %>%
select(Patient_SSN, Referral, Referring_Physician, Cost)

finaltable

```

```

##      Patient_SSN      Referral Referring_Physician Cost
##  1:  100000001 Christopher Turk      John Dorian    0
##  2:  100000001 Christopher Turk      John Dorian    0
##  3:  100000004   Percival Cox   Christopher Turk    0
##  4:  100000004   Percival Cox   Christopher Turk    0
##  5:  100000004   Percival Cox   Christopher Turk    0
##  ---
## 288: 100000102      John Wen      John Dorian 1500
## 289: 100000103      John Wen      John Dorian 1500
## 290: 100000103      John Wen      John Dorian 1500
## 291: 100000103 Christopher Turk      John Dorian 1500
## 292: 100000103 Christopher Turk      John Dorian 1500

```

*#Get the aggregation in the number of shared patients and total shared patient's cost*

```

num_shared_patients_costs <- finaltable %>%
  group_by(Referring_Physician, Referral) %>%
  summarise(shared_billing_costs= sum(Cost), Shared_patients = n_distinct(Patient_SSN)) %>% arrange(Referring_Physician)

```

## 'summarise()' has grouped output by 'Referring\_Physician'. You can override using the '.groups' argument

*# Adjusted\_Affiliated with table grouping by physician*

```

adjusted_affiliated_with <- affiliated_with %>%
  group_by(Physician) %>%
  filter(row_number() == 1) %>%
  select(Physician, Department)

```

*# Making Department Code-Department as referential keys*

```

adjusted_department_summary <- merge(adjusted_affiliated_with, department, by.x = 'Department', by.y='Department')

```

*# Making EmployeeID-Physician name as referential keys*

```

adjusted_physician_summary <- merge(adjusted_affiliated_with, physician, by.x='Physician', by.y='EmployeeID')
  rename(Physician_Name = Name) %>%
  select(Physician, Department, Physician_Name)

```

*# Combining the both adjusted summaries together*

```

adjusted_physician_department_summary <- merge(adjusted_physician_summary, adjusted_department_summary)
adjusted_physician_department_summary <- adjusted_physician_department_summary %>%
  select(Physician_Name, Department_Name)

```

*#Combining the final tables*

```

final_merged_result <- merge(num_shared_patients_costs, adjusted_physician_department_summary, by.x='Referring_Physician', by.y='Physician_Name')
  rename(Primary_Department=Department_Name) %>%
  merge(adjusted_physician_department_summary, by.x='Referral', by.y='Physician_Name') %>% rename(Referral_Department=Department_Name)
  arrange(Referring_Physician, desc(Shared_patients)) %>%
  select(Referring_Physician, Referral, Primary_Department, Referral_Department, Shared_patients, shared_billing_costs)

```



```
final_merged_result
```

```
## Referring_Physician Referral Primary_Department Referral_Department
## 1 Christopher Turk John Dorian General Medicine General Medicine
## 2 Christopher Turk Elliot Reid General Medicine General Medicine
## 3 Christopher Turk Molly Clock General Medicine Psychiatry
## 4 Christopher Turk Percival Cox General Medicine General Medicine
## 5 Christopher Turk John Wen General Medicine General Medicine
## 6 Christopher Turk Todd Quinlan General Medicine Surgery
## 7 Elliot Reid John Dorian General Medicine General Medicine
## 8 Elliot Reid Christopher Turk General Medicine General Medicine
## 9 Elliot Reid Molly Clock General Medicine Psychiatry
## 10 Elliot Reid Percival Cox General Medicine General Medicine
## 11 Elliot Reid John Wen General Medicine General Medicine
## 12 Elliot Reid Todd Quinlan General Medicine Surgery
## 13 John Dorian Elliot Reid General Medicine General Medicine
## 14 John Dorian Percival Cox General Medicine General Medicine
## 15 John Dorian Molly Clock General Medicine Psychiatry
## 16 John Dorian Christopher Turk General Medicine General Medicine
## 17 John Dorian John Wen General Medicine General Medicine
## 18 John Dorian Todd Quinlan General Medicine Surgery
## Shared_patients shared_billing_costs
## 1 17 0
## 2 15 0
## 3 9 0
## 4 8 0
## 5 4 33000
## 6 2 17298
## 7 18 0
## 8 9 31500
## 9 9 0
## 10 4 0
## 11 2 15000
## 12 2 7500
## 13 16 0
## 14 12 0
## 15 10 0
## 16 8 27700
## 17 7 30050
## 18 1 3750
```

### Question Three

Which tool did you find it 'easiest' to use while completing these exercises? What advice would you give novice data wranglers when it comes to choosing between Excel, SQL, and R? Please make your answer either a different text colour, or bolded, when you knit this document so TA's can find it.

**I found R to be the easiest while completing these exercises because R has a lot of built in functions/ packages that are necessary for me to clean data like removing unnecessary columns or null values. Excel has a lot of built in functions and I recommend to use the help function for all 3 platforms because the syntax can get messy when applying the same function to a specific column. When using Excel, I recommend to import some data and play around with the visual tools and some basic functions and queries necessary to transform and preprocess**

data. When using SQL, I recommend to be familiar with the common keywords that are used to query data. Once they get a grasp on the common keywords, then they can move on to some complex queries. Finally, when using R, the help function really should help novice data wranglers be familiar with the packages that are necessary to read in data, query data, subset dataframes, filter out data frames, and remove unnecessary column. They should also have a basic understanding on the functions used to generate common visual plots as well.

## Part Two

We are going to download US Census data using the Census API. To start, you will need to request a key here: [https://api.census.gov/data/key\\_signup.html](https://api.census.gov/data/key_signup.html).

We'll be using the following package:

A vignette demonstrating much of the functionality of this package can be found here <https://walker-data.com/census-r/index.html>

Start by setting your API key.

```
census_api_key("5c4e75b1d344c195de1c421444bb52400f92c18e")
```

```
## To install your API key for use in future sessions, run this function with 'install = TRUE'.
```

The function 'get\_acs()' will download the American Community Survey (ACS) Census data. You will need to know the variable ID - and there are thousands of variables across the different files. To rapidly search for variables, use the commands 'load\_variables()' and 'View()'. We'll do this below:

```
v19 <- load_variables(2019, "acs5", cache = TRUE)
View(v19)
```

As you can see, there are many types of data available to us in the census. In the View table, you can use filters to explore the kind of data that is available to you. For instance, try filtering by 'income' in the concept column.

The full metadata is available here [https://www.socialexplorer.com/data/ACS2019\\_5yr/metadata/](https://www.socialexplorer.com/data/ACS2019_5yr/metadata/).

For now, we'll use the following:

```
newEngDat <- get_acs(geography = "new england city and town area",
  year = 2019,
  variables = c(popn = "B03002_001",
    white = "B03002_003", blk = "B03002_004",
    asn = "B03002_006", hisp = "B03002_012",
    medHouseInc="B19013_001", hlthInsCov="B27001_001",
    workPop="B08604_001", workTravel="B08013_001",
    workHome="B08006_017", mthExp="B25088_001",
    mthHousing="B25105_001"),
  survey = "acs5",
  output = "wide")
```

```
## Getting data from the 2015-2019 5-year ACS
```

In the above code, we specified the following arguments:

**geography:** The level of geography we want the data in **year:** The end year of the data (because we want 2015-2019, we use 2019). **variables:** The variables we want to bring in as specified in a vector you create using the function `c()`. Note that we created variable names of our own (e.g. “popn”) and we put the ACS IDs in quotes (“B03002\_001”). **survey:** The specific Census survey we were extracting data from. We want data from the 5-year American Community Survey, so we specify “acs5”. The ACS comes in 1-, 3-, and 5-year varieties. **output:** gives us a traditional dataset, alternatively “tidy” would give us a tibble.

See `?get_acs` for more variables you could request.

We then have the following columns in our data:

**GEOID:** A unique ID variable of the geography **Name:** The Name of the geographic area **popn:** The total population **white:** The population of people who identify as white **blk:** The population of people who identify as black **asn:** The population of people who identify as asian **hisp:** The population of people who identify as hispanic **medHouseInc:** The median household income **hlthInsCov:** The population who have health insurance coverage **workPop:** The worker population **workTravel:** Aggregate travel time to work, in minutes **workHome:** Number of workers who work from home **mthExp:** Median monthly cost of living estimate **mthHousing:** Median housing costs per month

You’ll notice that there is an ‘E’ and an ‘M’ beside each of the column names in your dataset. The ‘E’ stands for estimate, and ‘M’ margin of error. While important, we will not be analyzing margins of error.

## Question Four

Remove the margin of error columns, and then remove the ‘E’ from the end of the other column names.

```
# your code here
```

```
# Cols stores the columns with the M removed  
cols <- c("GEOID", "NAME")
```

```
# This will store the updated columns with the Capital E removed at the end  
revised_col <- c("GEOID", "NAME")  
names_new_eng_dat <- names(newEngDat)
```

```
# I loop through starting the 3rd column and then first append the columns with E to a temporary vector
```

```
# cols and then I remove the "E, the last character of each column  
for (i in 3: length(names_new_eng_dat)){  
  if(str_sub(names_new_eng_dat[i], -1 ) == 'E'){  
    cols <- append(cols, names_new_eng_dat[i])  
    new_columns <- substr(names_new_eng_dat[i], 1, nchar(names_new_eng_dat[i]) - 1)  
    revised_col <- append(revised_col, new_columns)  
  }  
}
```

```
revised_col
```

```
## [1] "GEOID"      "NAME"      "popn"      "white"     "blk"  
## [6] "asn"        "hisp"      "medHouseInc" "hlthInsCov" "workPop"  
## [11] "workTravel" "workHome" "mthExp"     "mthHousing"
```

```
# Store the updated columns with 'M' removed at the end
newEngDat<- newEngDat[cols]
```

```
names(newEngDat) <- revised_col
newEngDat
```

```
## # A tibble: 40 x 14
##   GEOID NAME      popn white   blk   asn   hisp medHouseInc hlthInsCov workPop
##   <chr> <chr>   <dbl> <dbl> <dbl> <dbl> <dbl>      <dbl>      <dbl>   <dbl>
## 1 70450 Atho~ 2.15e4 1.99e4 118   122   841      54442      21367      7201
## 2 70600 Augu~ 7.70e4 7.29e4 699   547  1144      56876      75655     42372
## 3 70750 Bang~ 1.33e5 1.24e5 1186  1446  1704      54158     131960     66185
## 4 70900 Barn~ 2.41e5 2.14e5 6658  3355  7657      73714     238685    111832
## 5 71050 Barr~ 4.31e4 4.07e4 433   277   652      60882      42649     26168
## 6 71350 Benn~ 2.23e4 2.09e4 137   140   504      54797      21844     11380
## 7 71500 Berl~ 1.53e4 1.35e4 640    86   745      45597      13418      6195
## 8 71650 Bost~ 4.95e6 3.48e6 374214 390064 558137      91213     4905130    2794666
## 9 71950 Brid~ 9.21e5 5.83e5 97391 48097 170245      94222     914158     450822
## 10 73050 Dove~ 1.41e5 1.28e5 1254  4212  3617      72636     140114     58321
## # ... with 30 more rows, and 4 more variables: workTravel <dbl>,
## #   workHome <dbl>, mthExp <dbl>, mthHousing <dbl>
```

## Question 5

Which 10 communities have the largest proportion of their working population work from home?

```
# your code here
```

```
newEngDat2 <- newEngDat %>%
  mutate(workprop = workHome/workPop) %>%
  select(NAME, workprop) %>%
  arrange(desc(workprop))
head(newEngDat2,10)
```

```
## # A tibble: 10 x 2
##   NAME                                workprop
##   <chr>                                <dbl>
## 1 Vineyard Haven, MA Micropolitan NECTA 0.118
## 2 Keene, NH Micropolitan NECTA          0.0831
## 3 Greenfield Town, MA Micropolitan NECTA 0.0824
## 4 Barnstable Town, MA Metropolitan NECTA 0.0744
## 5 Bennington, VT Micropolitan NECTA      0.0731
## 6 Sanford, ME Micropolitan NECTA         0.0700
## 7 Brunswick, ME Micropolitan NECTA       0.0676
## 8 Portland-South Portland, ME Metropolitan NECTA 0.0671
## 9 Leominster-Gardner, MA Metropolitan NECTA 0.0656
## 10 Worcester, MA-CT Metropolitan NECTA    0.0631
```

## Question 6

We'll define discretionary income as Income-Expenses. Right now, you have annual income and monthly expenses. Create a new column to calculate monthly discretionary income, and display the towns with the

highest amounts of discretionary income.

```
# your code here
newEngDat$yearInc <- newEngDat$medHouseInc / 12
newEngDat$monthly_disc_in <- newEngDat$yearInc - newEngDat$smthExp
newEngDat_new <- newEngDat %>%
  select(NAME, "Discretionary_Income" = monthly_disc_in) %>%
  arrange(desc(Discretionary_Income))
head(newEngDat_new, 10)
```

```
## # A tibble: 10 x 2
##   NAME                               Discretionary_Income
##   <chr>                               <dbl>
## 1 Portsmouth, NH-ME Metropolitan NECTA 5648.
## 2 Danbury, CT Metropolitan NECTA      5644.
## 3 Bridgeport-Stamford-Norwalk, CT Metropolitan NECTA 5619.
## 4 Boston-Cambridge-Newton, MA-NH Metropolitan NECTA 5598.
## 5 Concord, NH Micropolitan NECTA      4868.
## 6 Hartford-East Hartford-Middletown, CT Metropolitan NECTA 4788.
## 7 Barnstable Town, MA Metropolitan NECTA 4661.
## 8 Manchester, NH Metropolitan NECTA    4650.
## 9 Lebanon, NH-VT Micropolitan NECTA    4631.
## 10 Worcester, MA-CT Metropolitan NECTA 4573.
```

### Question 7

Which 5 towns have the largest proportional gaps in healthcare coverage?

```
# your code here
newEngDat$prop_gap <- (newEngDat$popn - newEngDat$hlthInsCov) / newEngDat$popn
newEngDatprop <- newEngDat %>%
  select("Town Names" = NAME, "Proportional_Gap" = prop_gap) %>%
  arrange(desc(Proportional_Gap))
head(newEngDatprop, 5)
```

```
## # A tibble: 5 x 2
##   'Town Names'                     Proportional_Gap
##   <chr>                               <dbl>
## 1 Berlin, NH Micropolitan NECTA      0.120
## 2 Norwich-New London-Westerly, CT-RI Metropolitan NECTA 0.0443
## 3 Concord, NH Micropolitan NECTA     0.0372
## 4 Claremont, NH Micropolitan NECTA   0.0236
## 5 Bennington, VT Micropolitan NECTA  0.0201
```

### Question 8

The diversity index of a geographic area is the probability that two people selected at random will be the same race. Create a function which will sample from the reported ethnic population in each geographic area and return the diversity index. Display the top 5 diverse towns.

*# your code here*

```
diversity = function(area){
  newEngDat$whitepropind <- (newEngDat$white / newEngDat$popn) ** 2
  newEngDat$blackpropind <- (newEngDat$blk / newEngDat$popn) **2
  newEngDat$asianpropind <- (newEngDat$asn / newEngDat$popn) **2
  newEngDat$hispanicpropind <- (newEngDat$hispan / newEngDat$popn) **2
  total_index <- newEngDat$whitepropind + newEngDat$blackpropind+ newEngDat$asianpropind + newEngDat$hispanicpropind
  return(total_index)
}

head(diversity(newEngDat$NAME),5)
```

```
## [1] 0.8572081 0.8985311 0.8683063 0.7891973 0.8947059
```

### Question 9

Convert the ethnicity columns to be percentages. Make a boxplot where each ethnicity is represented on the x-axis, and percent is on the y-axis. Points will be awarded for ‘prettier’ plots!

*#your code here*

```
newEngDat$whiteprop <- (newEngDat$white / newEngDat$popn) * 100
newEngDat$blackprop <- (newEngDat$blk / newEngDat$popn) * 100
newEngDat$asianprop <- (newEngDat$asn / newEngDat$popn) * 100
newEngDat$hispanicprop <- (newEngDat$hispan / newEngDat$popn) * 100
newEngDat_Eth <- newEngDat %>%
  select(whiteprop, blackprop, asianprop, hispanicprop)
```

```
library(ggplot2)
library(lattice)
require(reshape2)
```

```
## Loading required package: reshape2
```

```
##
```

```
## Attaching package: 'reshape2'
```

```
## The following objects are masked from 'package:data.table':
```

```
##
```

```
##      dcast, melt
```

```
## The following object is masked from 'package:tidyr':
```

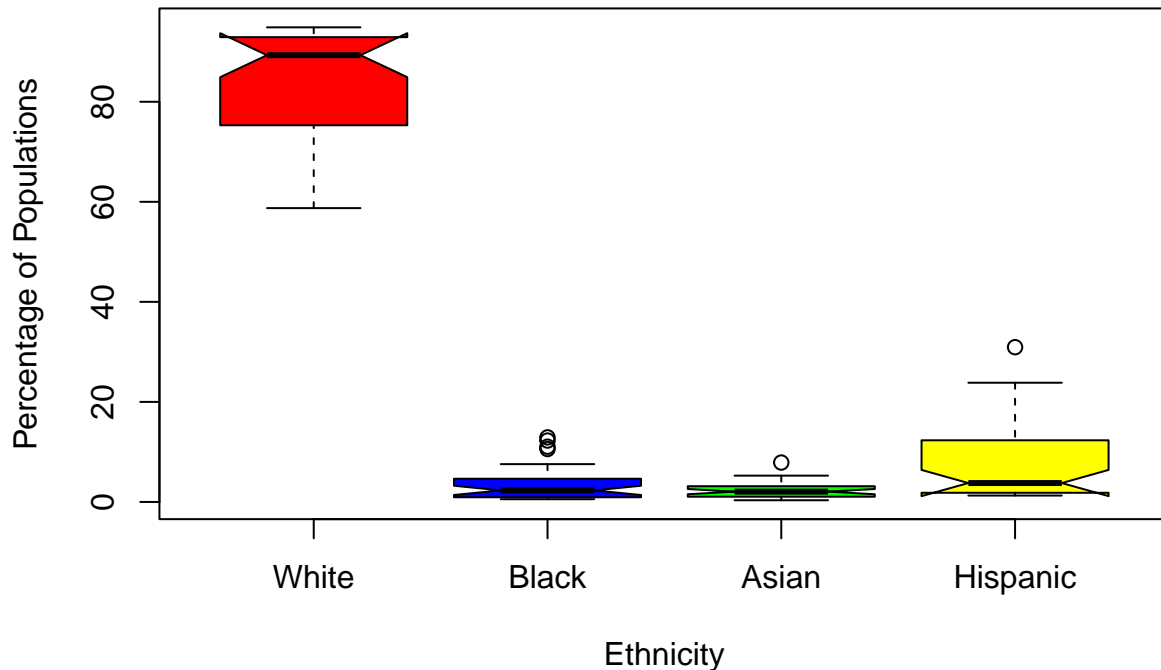
```
##
```

```
##      smiths
```

```
boxplot(newEngDat$whiteprop, newEngDat$blackprop, newEngDat$asianprop, newEngDat$hispanicprop, xlab = 'Ethnicity')
```

```
## Warning in (function (z, notch = FALSE, width = NULL, varwidth = FALSE, : some
## notches went outside hinges ('box'): maybe set notch=FALSE
```

## Relationship the proportion between all ethnic groups



### Question 10

Ask a question of your choosing. Output both the head of a table, and a simple plot answering your question. Feel free to use the API to import extra variables that may be of interest to you.

**Question:** Display the top 10 names with the highest poverty rate and make a simple plot establishing the relationship between poverty proportion and working proportion where poverty proportion is on the x axis and working proportion is on the Y axis

```
# I only select clumms that are necessary to answer my own question. I did not remove the E at the end
newEngDataq10 <- get_acs(geography = "new england city and town area",
  year = 2019,
  variables = c(popn = "B03002_001",
    pov = "B17023_001",
    hlthInsCov="B27001_001",
    workPop="B08604_001"),
  survey = "acs5",
  output = "wide")
```

```
## Getting data from the 2015-2019 5-year ACS
```

```
newEngDataq10
```

```
## # A tibble: 40 x 10
```

```
##      GEOID NAME      popnE popnM   povE   povM hlthInsCovE hlthInsCovM workPopE
##      <chr> <chr>      <dbl> <dbl>   <dbl> <dbl>      <dbl>      <dbl>      <dbl>
##  1 70450 Athol, MA M~ 2.15e4   233 5.46e3   257      21367        240       7201
##  2 70600 Augusta, ME~ 7.70e4   455 2.01e4   580      75655        494      42372
##  3 70750 Bangor, ME ~ 1.33e5   424 3.28e4   725     131960        455     66185
##  4 70900 Barnstable ~ 2.41e5    36 6.65e4  1114     238685        327    111832
##  5 71050 Barre, VT M~ 4.31e4   311 1.13e4   370      42649        367     26168
##  6 71350 Bennington,~ 2.23e4    75 5.30e3   261      21844        195     11380
##  7 71500 Berlin, NH ~ 1.53e4   131 3.80e3   186      13418        280      6195
##  8 71650 Boston-Camb~ 4.95e6   384 1.20e6  5068     4905130       1209   2794666
##  9 71950 Bridgeport~- 9.21e5   195 2.32e5  1932      914158        549   450822
## 10 73050 Dover-Durha~ 1.41e5    39 3.44e4   802     140114        295    58321
## # ... with 30 more rows, and 1 more variable: workPopM <dbl>
```

```
newEngDat_q10 <- newEngDataq10 %>%
  mutate(PovProp = povE/popnE)%>%
  mutate(workprop = workPopE / popnE) %>%
  select(NAME, PovProp, workprop) %>%
  arrange(desc(PovProp))

head(newEngDat_q10,10)
```

```
## # A tibble: 10 x 3
##      NAME                                     PovProp workprop
##      <chr>                                     <dbl>      <dbl>
##  1 Barnstable Town, MA Metropolitan NECTA      0.276      0.464
##  2 Portsmouth, NH-ME Metropolitan NECTA      0.269      0.721
##  3 Laconia, NH Micropolitan NECTA             0.269      0.503
##  4 Brunswick, ME Micropolitan NECTA            0.269      0.540
##  5 Barre, VT Micropolitan NECTA                0.263      0.608
##  6 Augusta, ME Micropolitan NECTA              0.262      0.551
##  7 Norwich-New London-Westerly, CT-RI Metropolitan NECTA 0.261      0.519
##  8 Pittsfield, MA Metropolitan NECTA           0.260      0.505
##  9 Lewiston-Auburn, ME Metropolitan NECTA      0.259      0.439
## 10 Rutland, VT Micropolitan NECTA              0.256      0.573
```

```
ggplot(newEngDat_q10 ,
  aes(x = PovProp,
      y = workprop)) +
  geom_point(color= "steelblue") +
  geom_smooth(method = "lm") + labs(x = "Poverty Proportion", y = "Working Proportion") + ggtitle("Rela
```

```
## 'geom_smooth()' using formula 'y ~ x'
```



Relationship Between Poverty Proportion and Working Proportion

